

Miguel Fuentes

Amherst, MA | 302-332-0176 | mmfuentes@umass.edu | miguel-fuentes.github.io | linkedin.com/in/miguel-fuentes-28679980 | github.com/Miguel-Fuentes

Summary

Ph.D. Candidate in Computer Science (graduating May 2026) specializing in differentially private synthetic data, probabilistic modeling, and machine learning. Demonstrated impact through publications in top venues (NeurIPS, AISTATS) and open-source releases. Seeking Research Scientist roles leveraging expertise in privacy, generative models, and applied ML for real-world challenges. Willing to relocate or work remotely.

Education

University of Massachusetts Amherst, PhD in Computer Science *Expected Graduation: May 2026*

- **Advisor:** Daniel Sheldon
- **Honors:** Spaulding-Smith Fellowship

University of Delaware, BS in Mathematics & BA in Computer Science *Aug 2016 – May 2020*

- **Honors:** University Honors Program General Honors Award

Skills

Technical Areas: Differential Privacy (Query Answering, Synthetic Data), Machine Learning, Generative Modeling, Graphical Models, Randomized Algorithms

Languages: Python, MATLAB, R

Libraries & Frameworks: JAX, PyTorch, NumPy, Matplotlib, Optax

Developer Tools: Git, Docker, SLURM, LaTeX, Snakemake, uv

Open Source: Contributed to github.com/ryan112358/private-pgm (10+ contributors, 100+ stars)

Research Experience

PhD Researcher, University of Massachusetts Amherst *Aug 2020 – Present*

- Developed a novel differentially private synthetic data algorithm that outperforms prior approaches (up to 40% error reduction); published at AISTATS 2024 and openly released. [github.com/Miguel-Fuentes/JAM_AiStats]
- Designed and implemented novel algorithms for efficient differentially private marginal query reconstruction, reducing workload error by $44\times$ over baseline. See code: github.com/bcmullins/efficient-marginal-reconstruction.
- Led foundational work on BirdFlow, a probabilistic model for bird migration using >1B citizen science records; project success enabled group to secure NSF funding and broader adoption in ecological research.
- Supported the development of a production-ready BirdFlow R package (as part of an NSF project), and maintain the core BirdFlow Python library to enable ongoing collaboration and reproducible science github.com/birdflow-science/BirdFlowPy.
- Built reproducible, large-scale experimental pipelines using SLURM, Snakemake, and Docker to streamline group research and accelerate development cycles.

Research Project Lead, Institute for Pure and Applied Mathematics *June 2019 – Aug 2019*

- Led a summer research project designing a privacy risk assessment framework for large-scale data analytics in Google Ads Data Hub, interfacing directly with a team of Google engineers.
- Authored a comprehensive technical report (43 pages) and presented findings both to Google Ads Data Hub leadership and at the Joint Mathematics Meetings (JMM) 2020.

University of Delaware Summer Scholars Program *June 2018 – Aug 2018*

- Investigated activation regions for shallow feed-forward neural networks, demonstrating (constructively) the existence of novel non-convex regions; results supported by theoretical proof and empirical validation.

Leadership Experience

Undergraduate Research Mentor

2023 – 2025

- Supervised an undergraduate student for two years on research applying implicit differentiation to Wasserstein distance optimization, resulting in an undergraduate thesis and acceptance into a PhD program.

EMBER Research Mentor

2022 – 2023

- Mentored four undergraduate students in a project on optimizing BirdFlow model performance. Guided experimental design and coding; students delivered a technical report and public symposium presentation.

Co-founder, UMass Differential Privacy Reading Group

2022 – 2024

- Initiated and organized a weekly reading group on foundational/contemporary topics in differential privacy, facilitating collaborations and ongoing research discussions.

Selected Publications

Efficient and Private Marginal Reconstruction with Local Non-Negativity

NeurIPS 2024

B. Mullins, **M. Fuentes** (co-first), Y. Xiao, D. Kifer, C. Musco, & D. Sheldon

<https://neurips.cc/virtual/2024/poster/93838>

Introduced GReM-LNN, a principled post-processing method that incorporates non-negativity constraints, reducing marginal query error by up to $44\times$ over baseline.

Joint Selection: Adaptively Incorporating Public Information for Private Synthetic Data

AISTATS 2024 (oral)

M. Fuentes (first author), B. Mullins, R. McKenna, G. Miklau, & D. Sheldon

<https://proceedings.mlr.press/v238/fuentes24a/fuentes24a.pdf>

Developed JAM-PGM, the first graphical-model-based mechanism to jointly use public/private sources; achieved up to 40% error reduction on benchmarks.

BirdFlow: Learning seasonal bird movements from eBird data

Methods in Ecology and Evolution, 2023

M. Fuentes (first author), B. Van Doren, D. Fink, & D. Sheldon

<https://doi.org/10.1111/2041-210X.14052>

Created a probabilistic model for large-scale migration forecasting; code and models used in Cornell Lab collaborations.

Full publication list: scholar.google.com/citations?user=wy3LBgsAAAAJ&hl

Presentations

- Oral: “Joint Selection: Adaptively Incorporating Public Information for Private Synthetic Data” at AISTATS 2024 (Valencia, Spain) and Theory and Practice of Differential Privacy (Boston, MA), 2024.